

WEIDI LUO

The University of Georgia, Herty Dr, Athens, GA 30602

luo.1455@uga.edu www.eddyluo.com

RESEARCH INTERESTS

Trustworthy AI & AI Safety and Security

EDUCATION

School of Computing, University of Georgia Aug 2025-Present

P.H.D. Student in Computer Science

- Advisor: Prof.Zhen Xiang and Prof.Chaowei Xiao

Art and Science College, The Ohio State University Aug 2022-Jul 2025

B.A. in Computer and Information Science

- Advisor: Prof.Chaowei Xiao, Dr.Yu Li, Prof.Ruixiang Tang, Prof.Yu Su

WORK EXPERIENCES

Research Scientist | Virtue AI | San Francisco, USA May 2026-Aug 2026

- Advisor: Wenbo Guo, Head of Agent Security

Large Language Model Engineer | IDEA | Shen Zhen, China Jul 2024-Oct 2024

- Advisor: Yu Li, Principal Researcher

PROFESSIONAL EXPERIENCES

Research Intern | Johns Hopkins University Dec 2023-Apr 2025

- Advisor: Chaowei Xiao, Assistant Professor at Department of Electrical and Computer Engineering

Research Intern | Rutgers University Oct 2024-Dec 2024

- Advisor: Ruixiang Tang, Assistant Professor at Department of Computer Science

Research Assistant | ICICLE Institute | The Ohio State University Apr 2023-Aug 2023

- Advisor: Yu Su, Distinguished Assistant Professor at College of Engineering

REVIEWER EXPERIENCES

ACL'2025, EMNLP'2025, EACL'2026, ACL'2026, ICML'2026, COLM'2026, NeurIPS'2026

HOBBIES

Violin (Level 10, Sichuan Musicians Association)

PUBLICATION

Weidi Luo*, Qiming Zhang*, Tianyu Lu*, Xiaogeng Liu, Yue Zhao, Jieyu Zhao, Song Gao, Patrick McDaniel, Zhen Xiang, Chaowei Xiao, “*Doxing via the Lens: Revealing Location-related Privacy Leakage on Multi-modal Large Reasoning Models*”, Accepted by ICLR'2026.

Weidi Luo, Shenghong Dai, Xiaogeng Liu, Suman Banerjee, Huan Sun, Muhao Chen, Chaowei Xiao, “*AGrail: A Lifelong Agent Guardrail with Effective and Adaptive Safety Detection*”, Accepted by ACL'2025 (Main).

Weidi Luo*, Siyuan Ma*, Xiaogeng Liu*, Xiaoyu Guo, Chaowei Xiao, “*JailBreakV-28K: A Benchmark for Assessing the Robustness of MultiModal Large Language Models against Jailbreak Attacks*”, Accepted by COLM'2024.

Weidi Luo*, He Cao*, Yu Wang, Zijing Liu, Aidan Wong, Bin Feng, Yuan Yao, Yu Li, “*Guide for Defense (G4D): Dynamic Guidance for Robust and Balanced Defense in Large Language Models.*”, Accepted by NAACL'2025 (Findings).

Mingyu Jin, **Weidi Luo**, Sitao Cheng, Xinyi Wang, Wenye Hua, Ruixiang Tang, William Yang Wang, Yongfeng Zhang “*Disentangling Memory and Reasoning Ability in Large Language Models*”, Accepted by ACL'2025 (Main).

Vardaan Pahuja, **Weidi Luo**, Yu Gu, Cheng-Hao Tu, Hong-You Chen, Tanya Berger-Wolf, Charles Stewart, Song Gao, Wei-Lun Chao, Yu Su, “*Bringing Back the Context: Image Classification as Link Prediction on Multimodal Knowledge Graphs*”, Accepted

by CIKM'2024 (Oral).

PRE-PRINT

Weidi Luo, Qiming Zhang, Tianyu Lu, Xiaogeng Liu, Bin Hu, Hung-Chun Chiu, Siyuan Ma, Yizhe Zhang, Xusheng Xiao, Yinzhi Cao, Zhen Xiang, Chaowei Xiao, “*Code Agent can be an End-to-end System Hacker: Benchmarking Real-world Threats of Computer-use Agent*”, on Arxiv

Weidi Luo*, Xiaofei Wen*, Tenghao Huang*, Hongyi Wang, Zhen Xiang, Chaowei Xiao, Kristina Gligorić, Muhao Chen, "Cooking Up Risks: Benchmarking and Reducing Food Safety Risks in Large Language Models", on Arxiv.

Siyuan Ma*, **Weidi Luo***, Yu Wang, Xiaogeng Liu, Muhao Chen, Bo Li, Chaowei Xiao, “*Visual-RolePlay: Universal Jailbreak Attack on MultiModal Large Language Models via Role-playing Image Character*”, on Arxiv.

Zeru Shi, Zhenting Wang, Yongyue Su, **Weidi Luo**, Fan Yang, Yongfeng Zhang, “*Robustness-aware Automatic Prompt Optimization*”, on Arxiv.

AWARD

- \$20,000 SafeBench Award from Center for AI Safety, 2025
- Graduate Student Research Award from University of Georgia , School of Computing, 2026