

✉ luo.1455@uga.edu
🏠 eddyluo.com
🔍 [google scholar](https://scholar.google.com/)
🐙 github.com/EddyLuo1232
🐦 x.com/luoweidi84
🌐 linkedin.com/in/weidi-luo

School of Computing
University of Georgia
Athens, GA
United States

Research Overview

My research focuses on **AI safety and security**, with a particular emphasis on trustworthy foundation models, including multimodal and reasoning models, and autonomous AI agents. Across my research and industry experience, I study how advanced models fail under realistic adversarial, privacy, and security threats, and how to build evaluations and safeguards that remain effective in deployment.

- **AI Safety and Security:** detecting real-world threats, jailbreak and privacy-risk analysis, mechanistic safety analysis, and risk mitigation for foundation models.
- **Agent Safety:** adaptive agent guardrails, agent interpretability, computer-use and coding-agent red-teaming, and real-world threat benchmarking.
- **AI for Security Applications:** vulnerability discovery, cybersecurity evaluation, and security-focused applications of foundation models and AI agents.

Education

Ph.D. Student CS Aug 2025 – Present	University of Georgia, School of Computing Advisor: Zhen Xiang Co-advised by Chaowei Xiao, Johns Hopkins University
B.A. CS Aug 2022 – Jul 2025	The Ohio State University, College of Arts and Sciences Advisors: Chaowei Xiao, Yu Li, Ruixiang Tang, and Yu Su

Industry Experience

Research Scientist Jul 2026 – Aug 2026	Meta Superintelligence Labs Menlo Park, CA, USA Advisor: Bo Li
Research Scientist May 2026 – Jun 2026	Virtue AI San Francisco, CA, USA Advisor: Zhaorun Chen
MLE Intern Jul 2024 – Oct 2024	IDEA Shenzhen, Guangdong, China Advisor: Yu Li

Research Experience

Research Intern Dec 2023 – Apr 2025	Johns Hopkins University Department of Electrical and Computer Engineering Advisor: Chaowei Xiao
Research Intern Oct 2024 – Dec 2024	Rutgers University Department of Computer Science Advisor: Ruixiang Tang
Research Assistant Apr 2023 – Aug 2023	ICICLE Institute, The Ohio State University College of Engineering Advisor: Yu Su

Honors and Awards

SoC Graduate Student Research Award, University of Georgia	2026
\$20,000 SafeBench Prize, Center for AI Safety	2025

Selected Publications

The * sign denotes equal contribution.

Peer-Reviewed Papers

- C7 Weidi Luo*, Qiming Zhang*, Tianyu Lu*, Xiaogeng Liu, Bin Hu, Hung-Chun Chiu, Siyuan Ma, Yizhe Zhang, Xusheng Xiao, Yinzhi Cao, Zhen Xiang, and Chaowei Xiao. **Your Harness is Not Secure: Benchmarking Real-world Threat of Command Line Interface Agent**. In *Empirical Methods in Natural Language Processing (EMNLP) Findings*. 2026.
- C6 Weidi Luo*, Tianyu Lu*, Qiming Zhang*, Xiaogeng Liu, Bin Hu, Yue Zhao, Jieyu Zhao, Song Gao, Patrick McDaniel, Zhen Xiang, and Chaowei Xiao. **Doxing via the Lens: Revealing Location-related Privacy Leakage on Multi-modal Large Reasoning Models**. In *International Conference on Learning Representations (ICLR)*. 2026.
- C5 Weidi Luo, Shenghong Dai, Xiaogeng Liu, Suman Banerjee, Huan Sun, Muhao Chen, and Chaowei Xiao. **AGrail: A Lifelong Agent Guardrail with Effective and Adaptive Safety Detection**. In *Annual Meeting of the Association for Computational Linguistics (ACL) Main*. 2025.
- C4 Weidi Luo*, Siyuan Ma*, Xiaogeng Liu*, Xiaoyu Guo, and Chaowei Xiao. **JailBreakV-28K: A Benchmark for Assessing the Robustness of MultiModal Large Language Models against Jailbreak Attacks**. In *Conference on Language Modeling (COLM)*. 2024. Winner of the \$20,000 SafeBench Prize; integrated into NVIDIA/garak and Microsoft's PyRIT.
- C3 Weidi Luo*, He Cao*, Yu Wang, Zijing Liu, Aidan Wong, Bin Feng, Yuan Yao, and Yu Li. **Dynamic Guided and Domain Applicable Safeguards for Enhanced Security in Large Language Models**. In *Annual Conference of the North American Chapter of the Association for Computational Linguistics (NAACL) Findings*. 2025.
- C2 Mingyu Jin, Weidi Luo, Sitao Cheng, Xinyi Wang, Wenyue Hua, Ruixiang Tang, William Yang Wang, and Yongfeng Zhang. **Disentangling Memory and Reasoning Ability in Large Language Models**. In *Annual Meeting of the Association for Computational Linguistics (ACL) Main*. 2025.
- C1 Vardaan Pahuja, Weidi Luo, Yu Gu, Cheng-Hao Tu, Hong-You Chen, Tanya Berger-Wolf, Charles Stewart, Song Gao, Wei-Lun Chao, and Yu Su. **Bringing Back the Context: Image Classification as Link Prediction on Multimodal Knowledge Graphs**. In *ACM International Conference on Information and Knowledge Management (CIKM) Oral*. 2024.

Preprints

P3 Weidi Luo, Qiming Zhang, Yihao Quan, Mingyu Jin, Jie Cai, Chaowei Xiao, Jingcheng Niu, and Zhen Xiang. **AgentLens: Interpretable Safety Steering via Mechanistic Subspaces for Multi-Turn Coding Agent**. In *arXiv Preprint*. 2026.

P2 Weidi Luo*, Xiaofei Wen*, Tenghao Huang*, Hongyi Wang, Zhen Xiang, Chaowei Xiao, Kristina Gligoric, and Muhao Chen. **Cooking Up Risks: Benchmarking and Reducing Food Safety Risks in Large Language Models**. In *arXiv Preprint*. 2026.

P1 Siyuan Ma*, Weidi Luo*, Yu Wang, Xiaogeng Liu, Muhao Chen, Bo Li, and Chaowei Xiao. **Visual-RolePlay: Universal Jailbreak Attack on MultiModal Large Language Models via Role-playing Image Character**. In *arXiv Preprint*. 2024.

Professional Service

Conference Reviewing

Annual Meeting of the Association for Computational Linguistics (ACL)	2025, 2026
Empirical Methods in Natural Language Processing (EMNLP)	2025
European Chapter of the Association for Computational Linguistics (EACL)	2026
International Conference on Machine Learning (ICML)	2026
Conference on Language Modeling (COLM)	2026
Neural Information Processing Systems (NeurIPS)	2026
Association for the Advancement of Artificial Intelligence (AAAI)	2027

Journal Reviewing

IEEE Transactions on Information Forensics and Security (IEEE T-IFS)	2026
Proceedings of the IEEE	2026

Hobbies

Violin, Level 10, Sichuan Musicians Association	
---	--